NWEN 242

Pipelined Datapath and Control

Agenda

- Single-cycle vs. pipelined performance
- Pipeline register and pipeline operation
- Control for the pipeline

Introduction to pipelining

- Involves having several instructions in the processor all at once
  - Each at a different stage of its execution
  - Individual instructions take the same time (five steps)

5-stage pipelined datapath

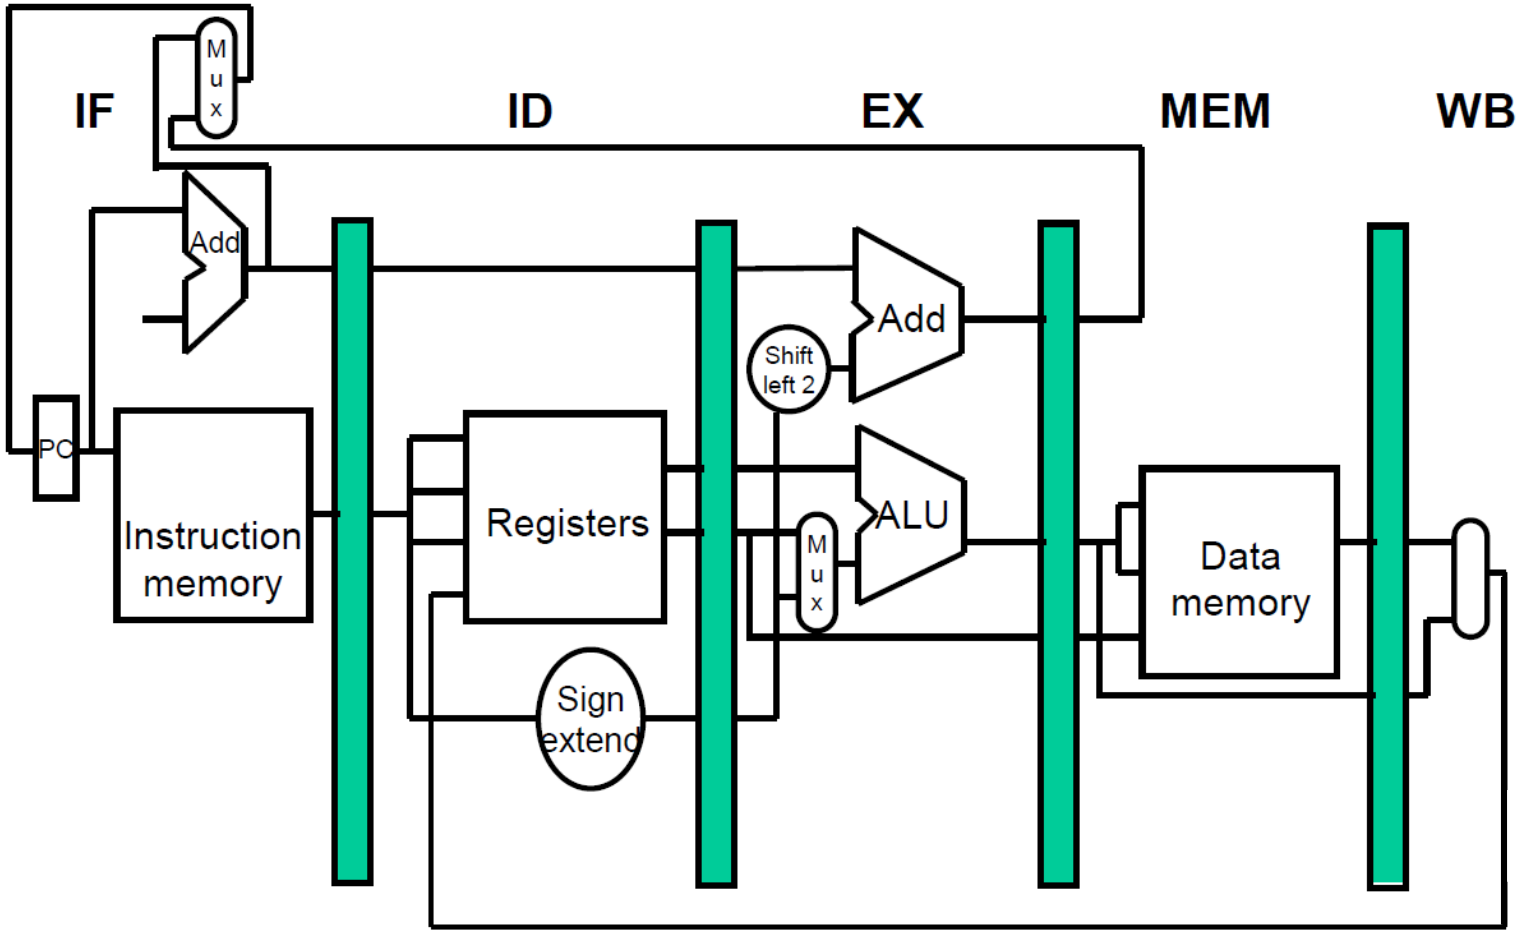

Benefits of pipelining

- Pipelining gives better performance than single-cycle design
- Improved throughput

Throughput for single-cycle datapath

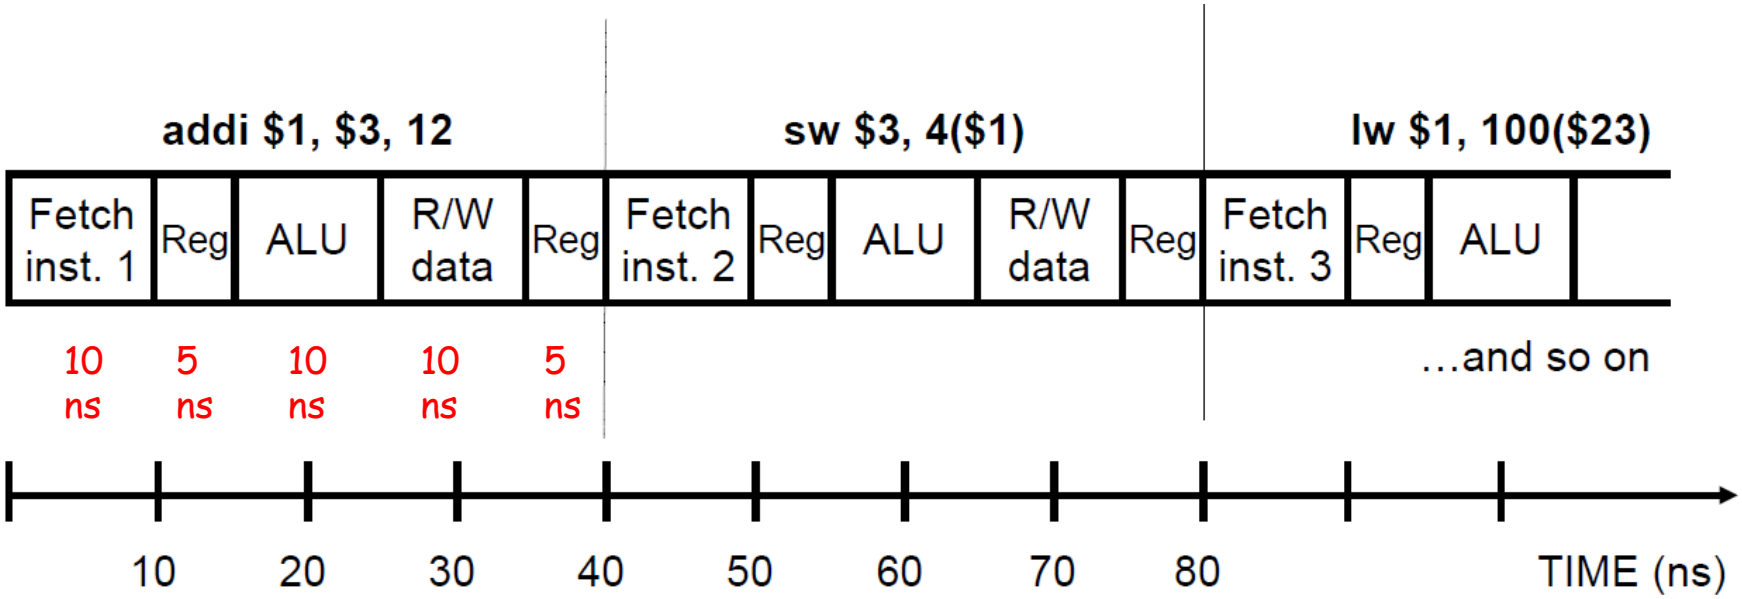

- Assumed latencies
  - Memory 10ns
  - ALU 10ns
  - Register 5ns
- Suitable clock speed
  - 25 MHz
  - 40ns period
- 120 ns to execute the three instructions

Throughput for pipelined datapath

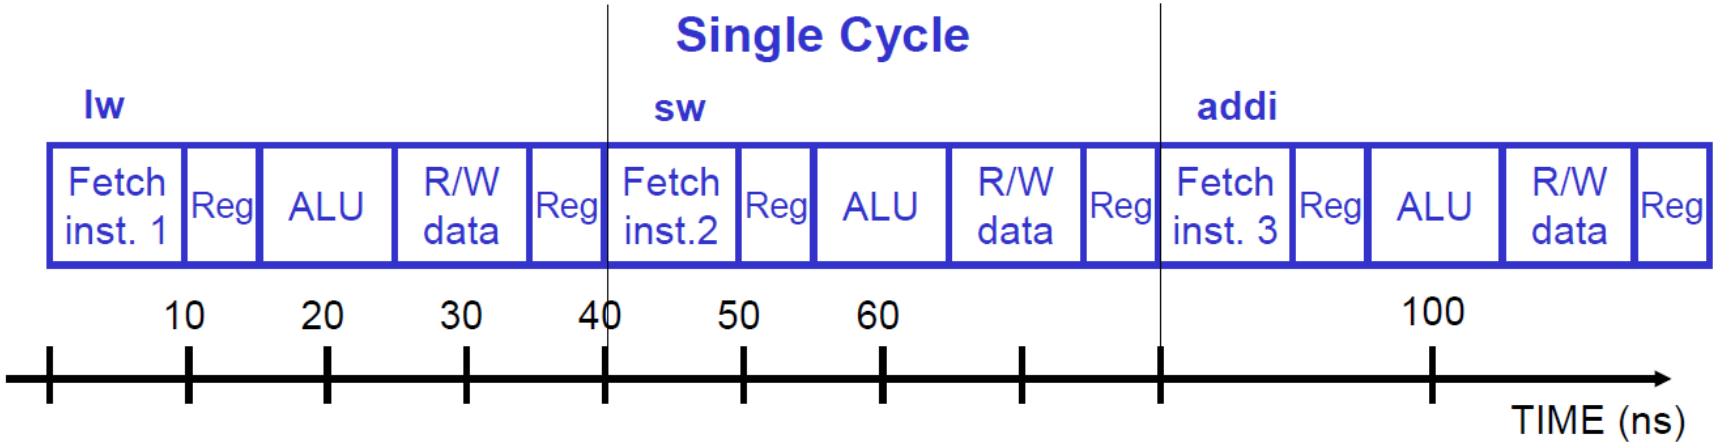

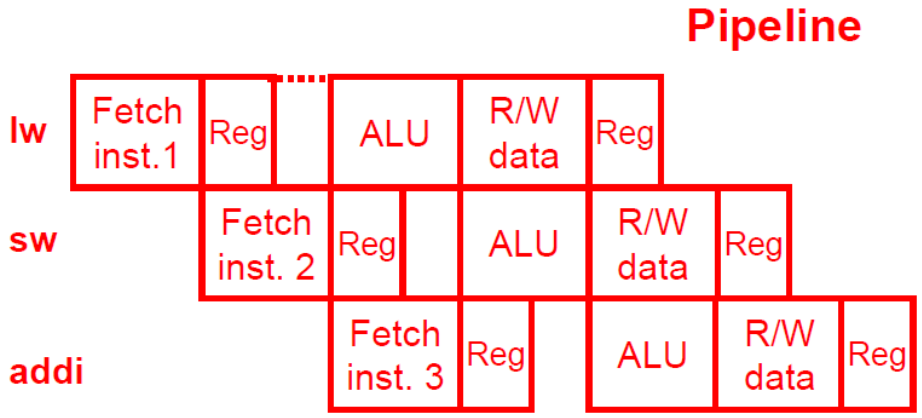

- Assumed latencies
  - Memory 10ns
  - ALU 10ns
  - Register 5ns
- 70 ns to execute the three instructions

Total clock cycles = 5 + (n-1)

What is the performance gain?

- Non-pipelined
  - Each instruction takes 40ns to complete
  - e.g. completing 3 instructions takes 120ns
- Pipelined with 5 stages
  - Each instruction takes 50ns to complete
  - e.g. third instruction finishes after 50+20=70ns

Time to complete the first instruction

Time to complete next two instructions

- Thus pipelined processor is 120/70 = 1.71 faster (for three instructions) than non-pipelined processor
- How many times faster is the pipelined processor for executing 1000 instructions?

- Single-cycle processor: (4 x 1000) stages x 10ns = 40,000ns
- Pipelined processor: (5 + (1000-1)) stages x 10ns = 10,040ns
- Pipelined processor is almost four times faster than non-pipelined counterpart.

- Single-cycle processor: (4 x 1000) stages x 10ns = 40,000ns
- Pipelined processor: (5 + (1000-1)) stages x 10ns = 10,040ns
- Pipelined processor is almost four times faster than non-pipelined counterpart.

Performance Gain = single-cycle cycle time / pipelined cycle time

How pipeline works

- New datapath element: Pipeline Registers

Pipeline registers

- Need registers between stages
  - to hold information produced in previous cycle
  - get written for every cycle

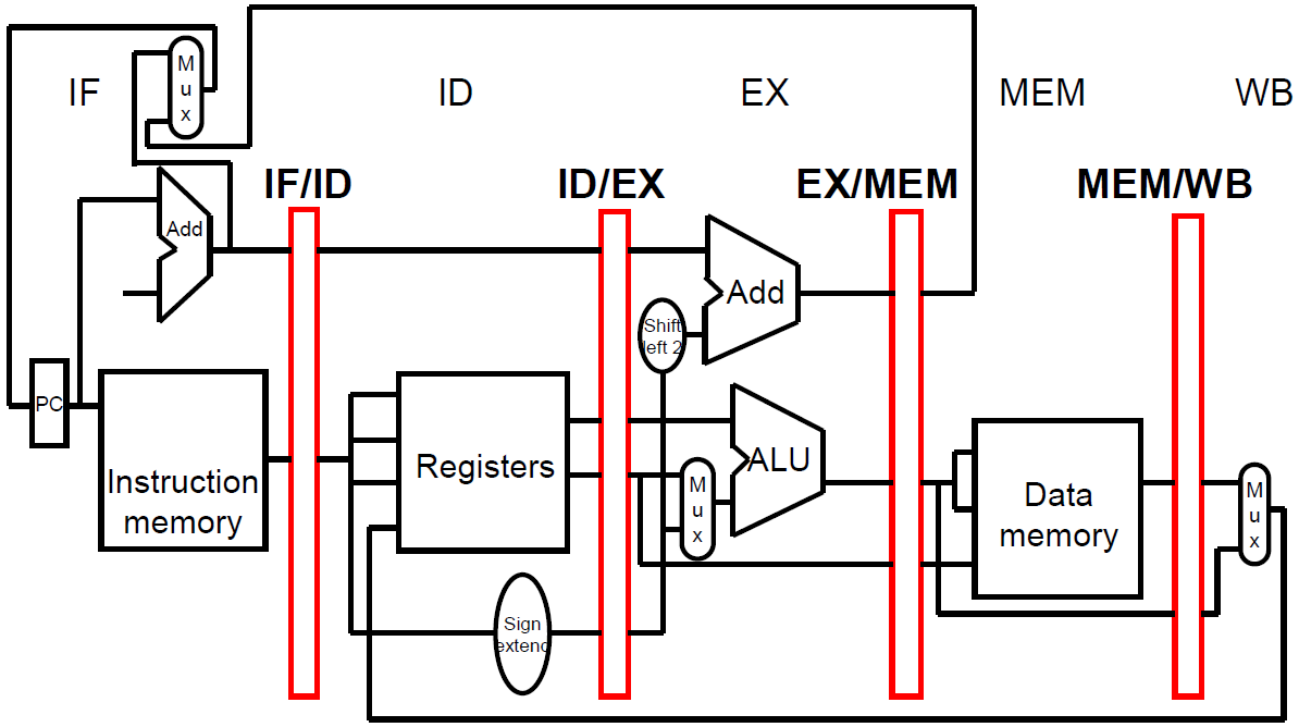

Pipeline operation

- Cycle-by-cycle flow of instructions through the pipelined datapath

lw stage 1 (IF = instruction fetch)

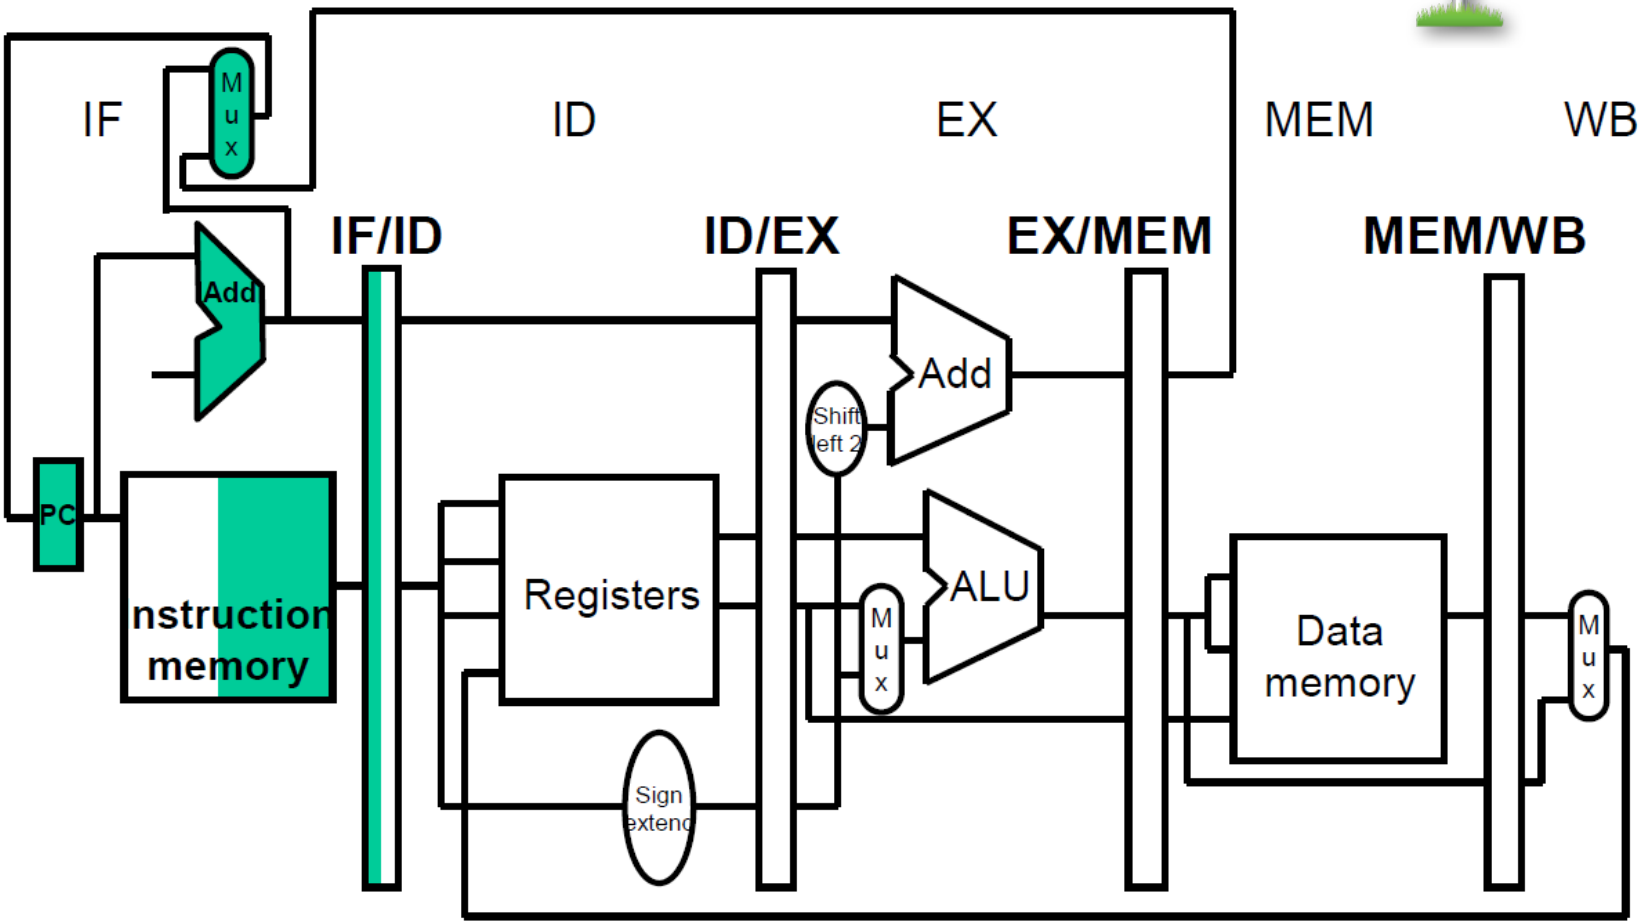

- Right half of a block colored means it's being read
- Left half of a block colored means it's being written

lw stage 2 (ID = instruction decode)

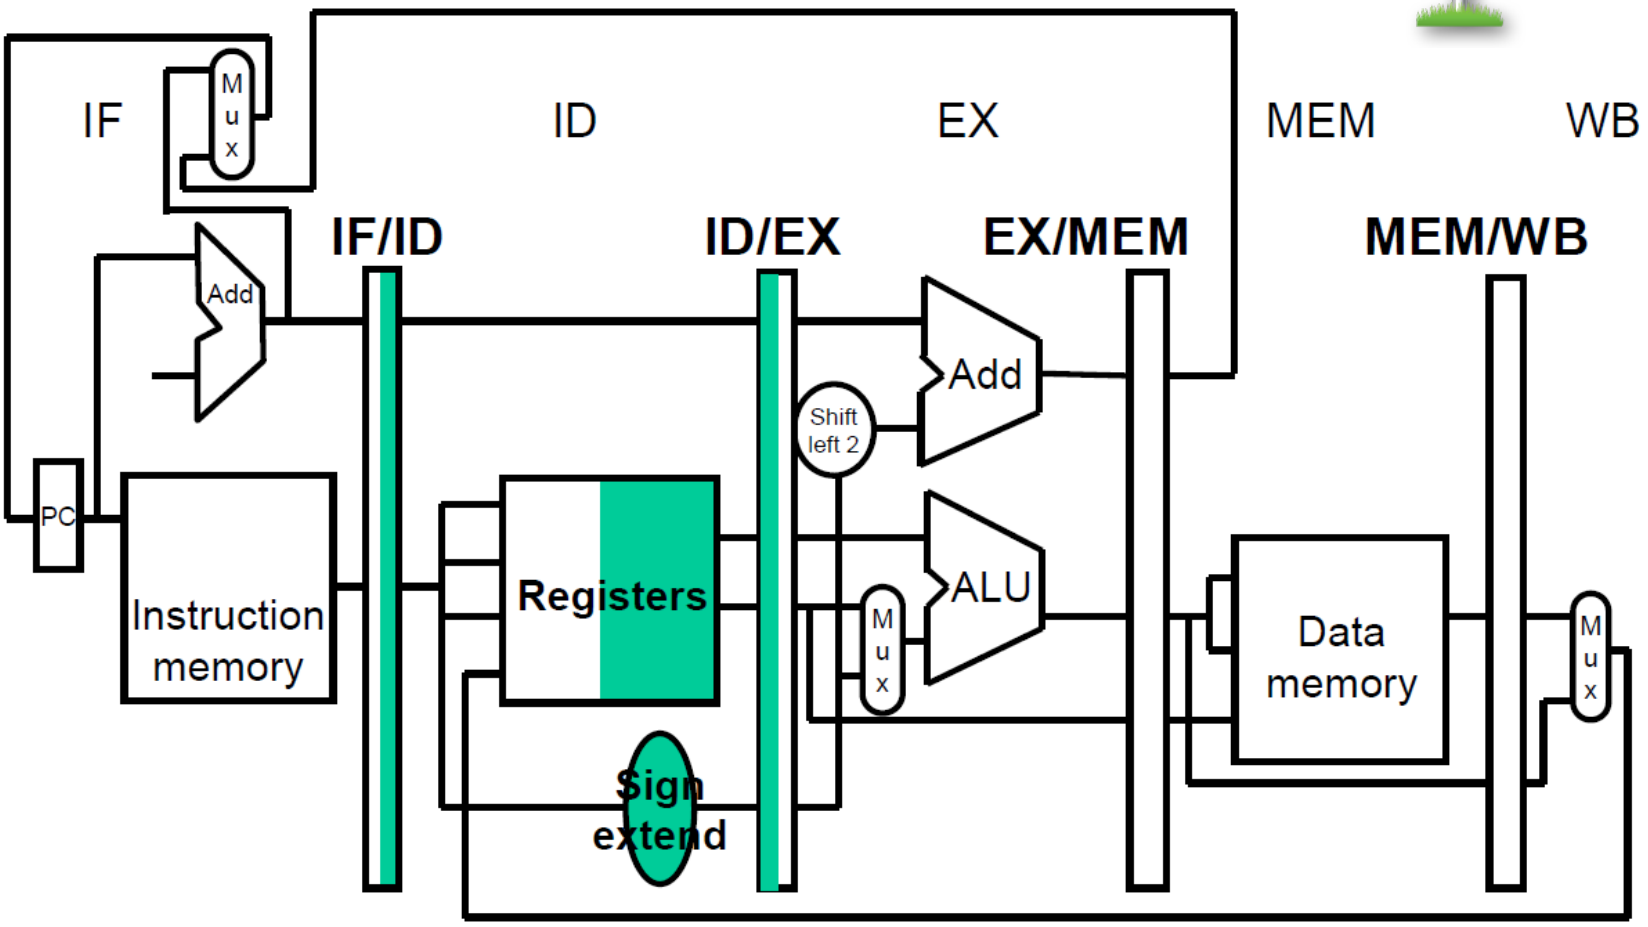

lw stage 3 (EX = execute)

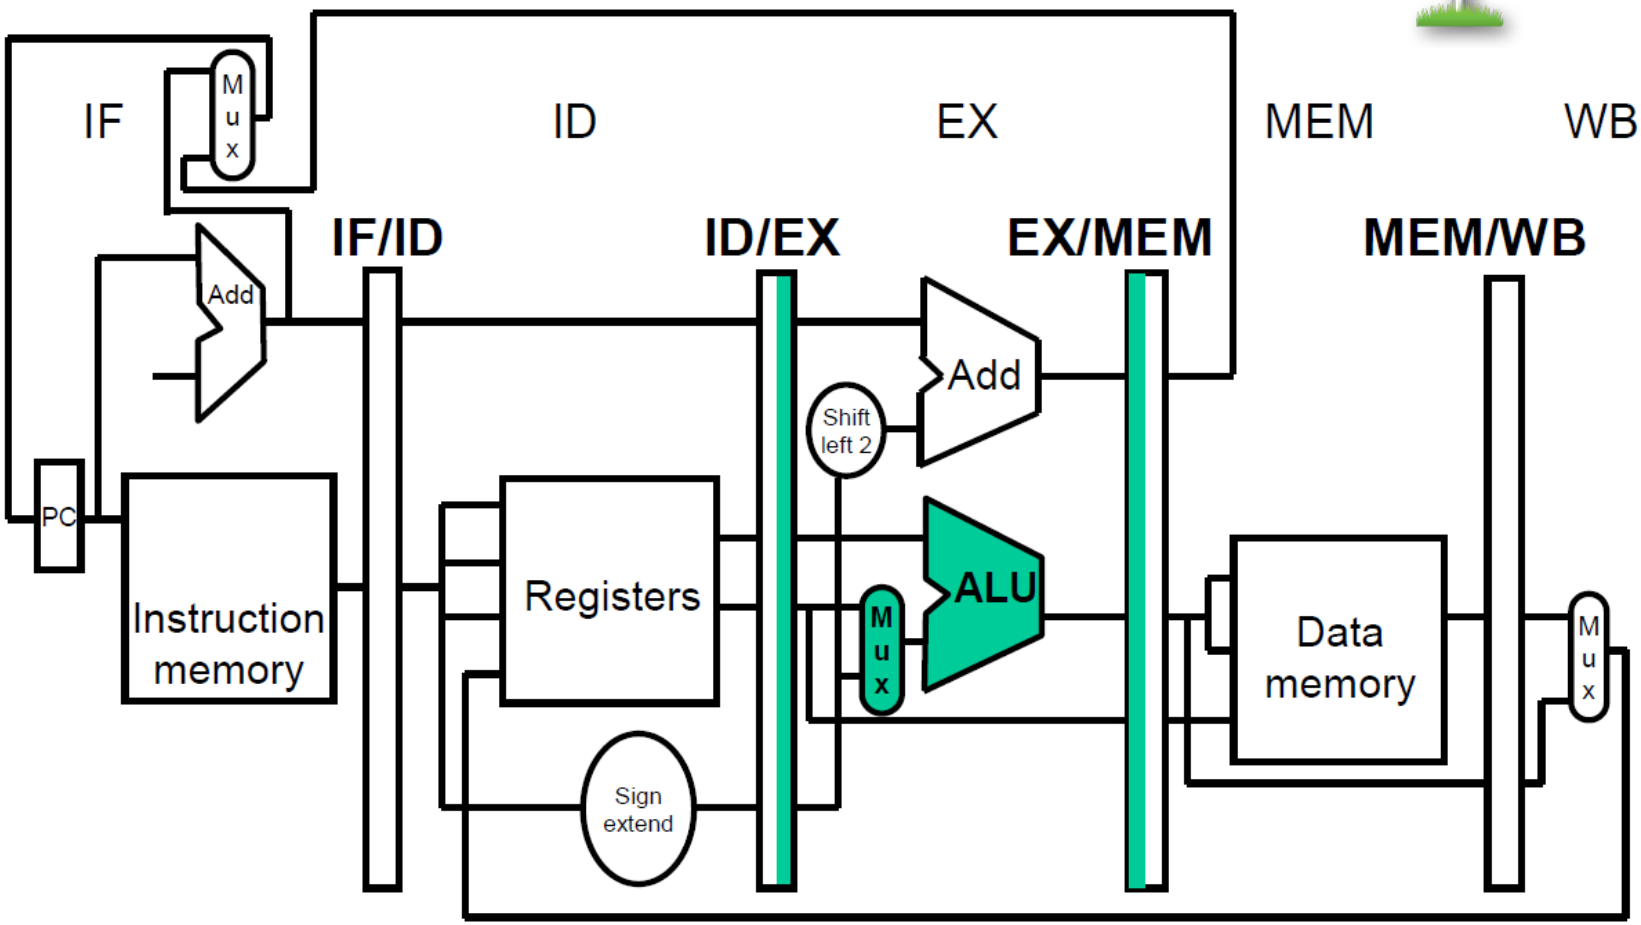

lw stage 4 (MEM = R/W memory)

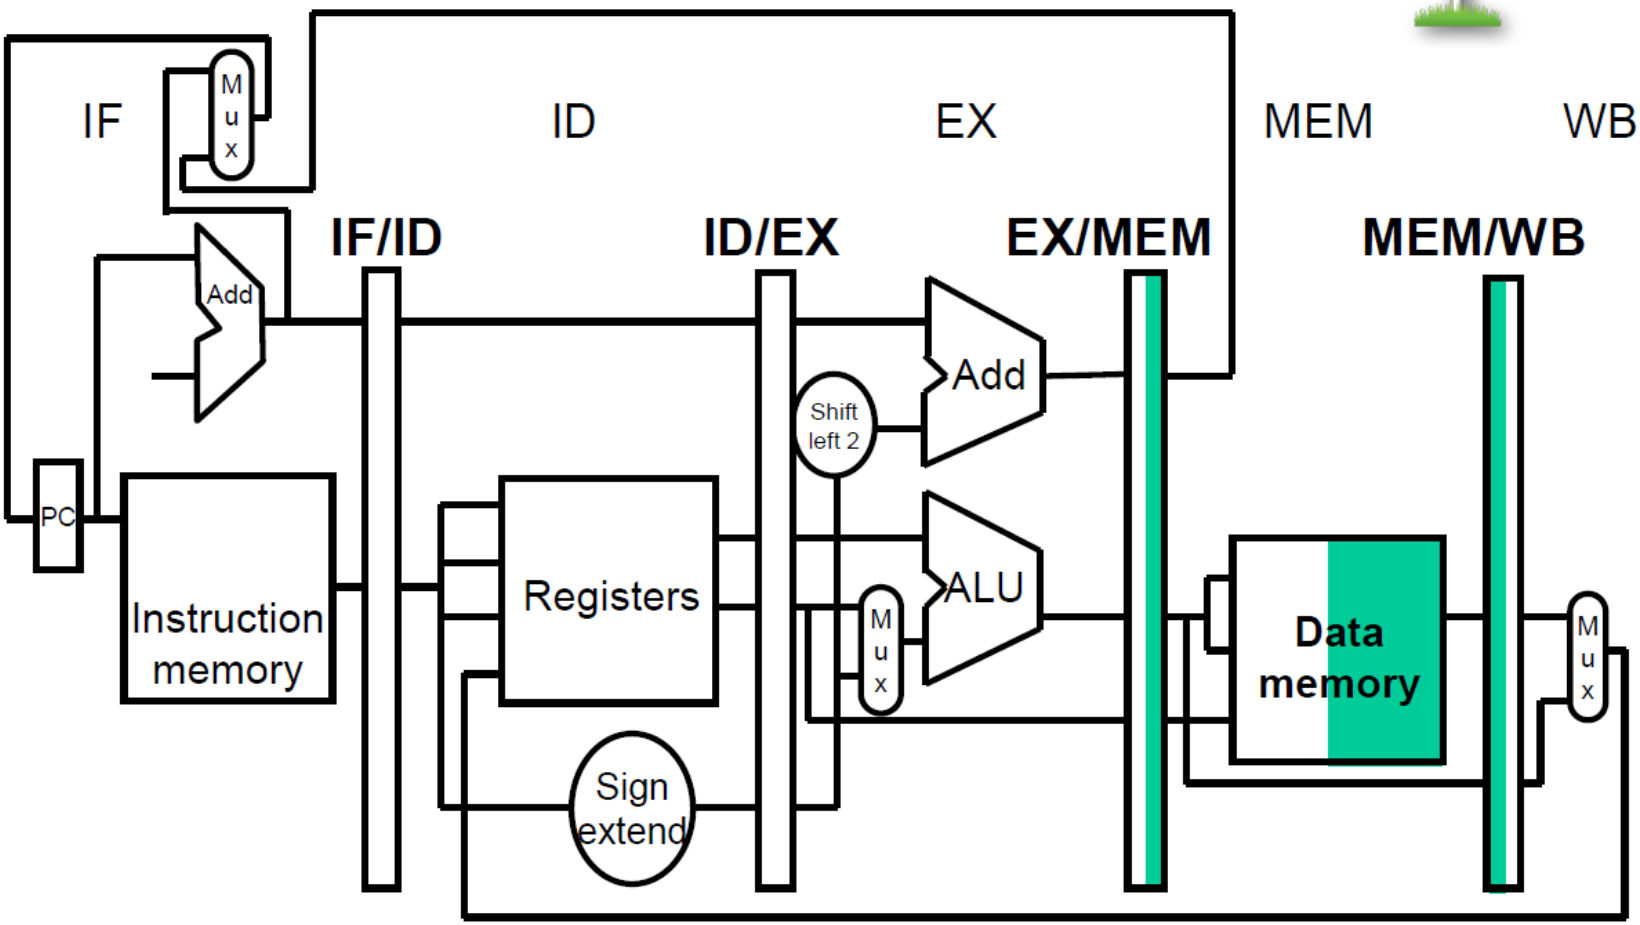

lw stage 5 (WB = Write-Back)

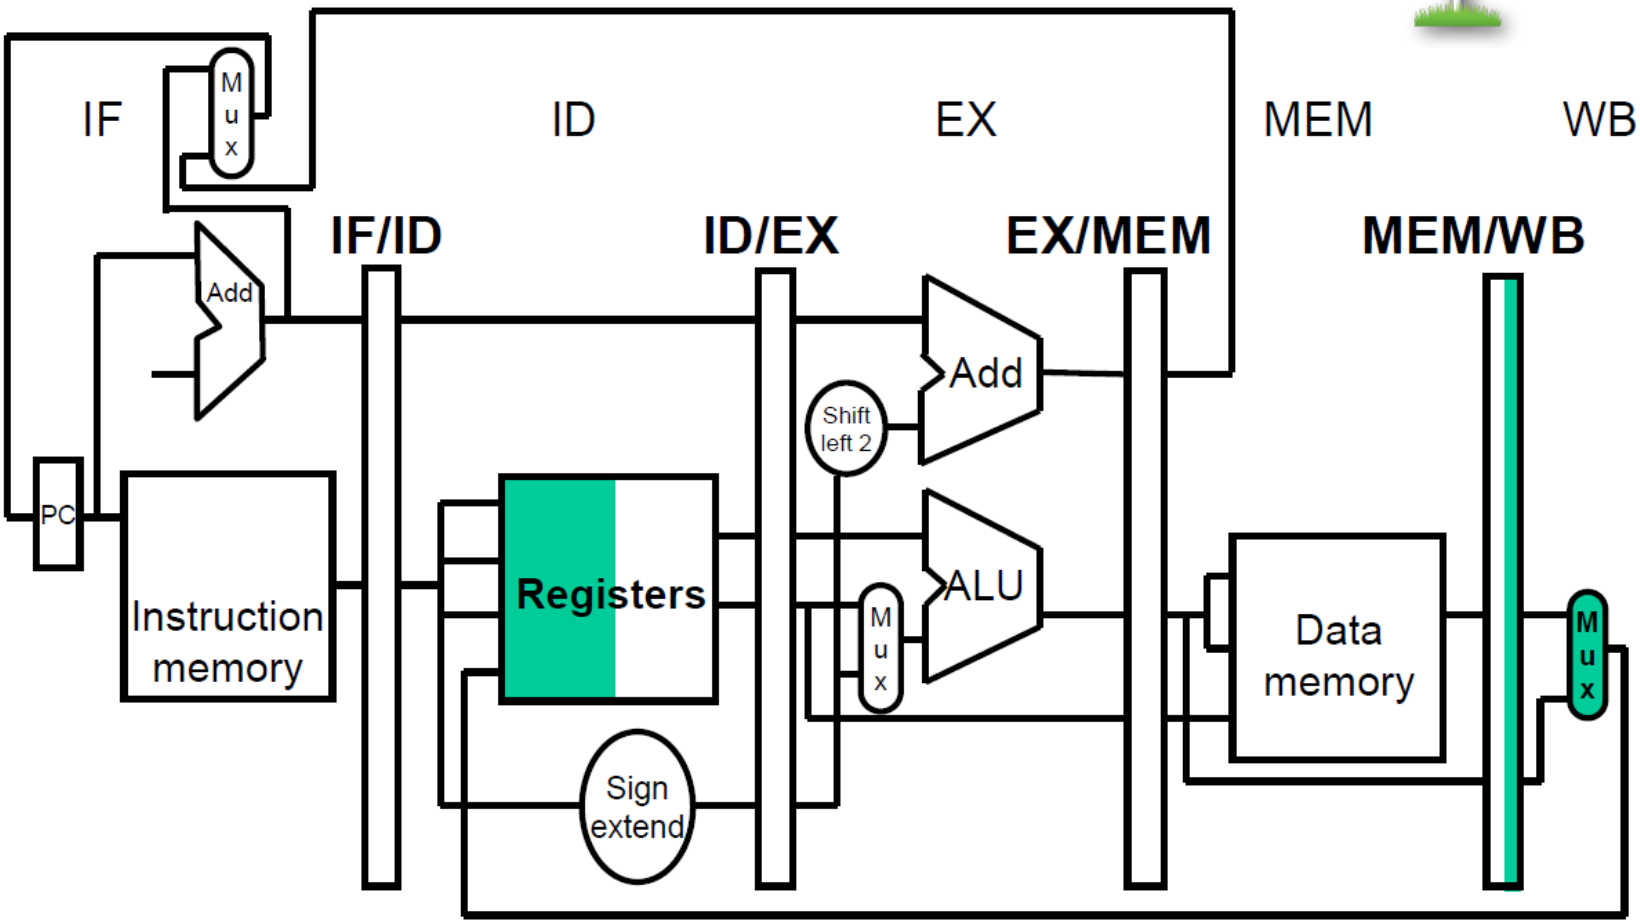

lw stage 5 (WB = Write-Back)

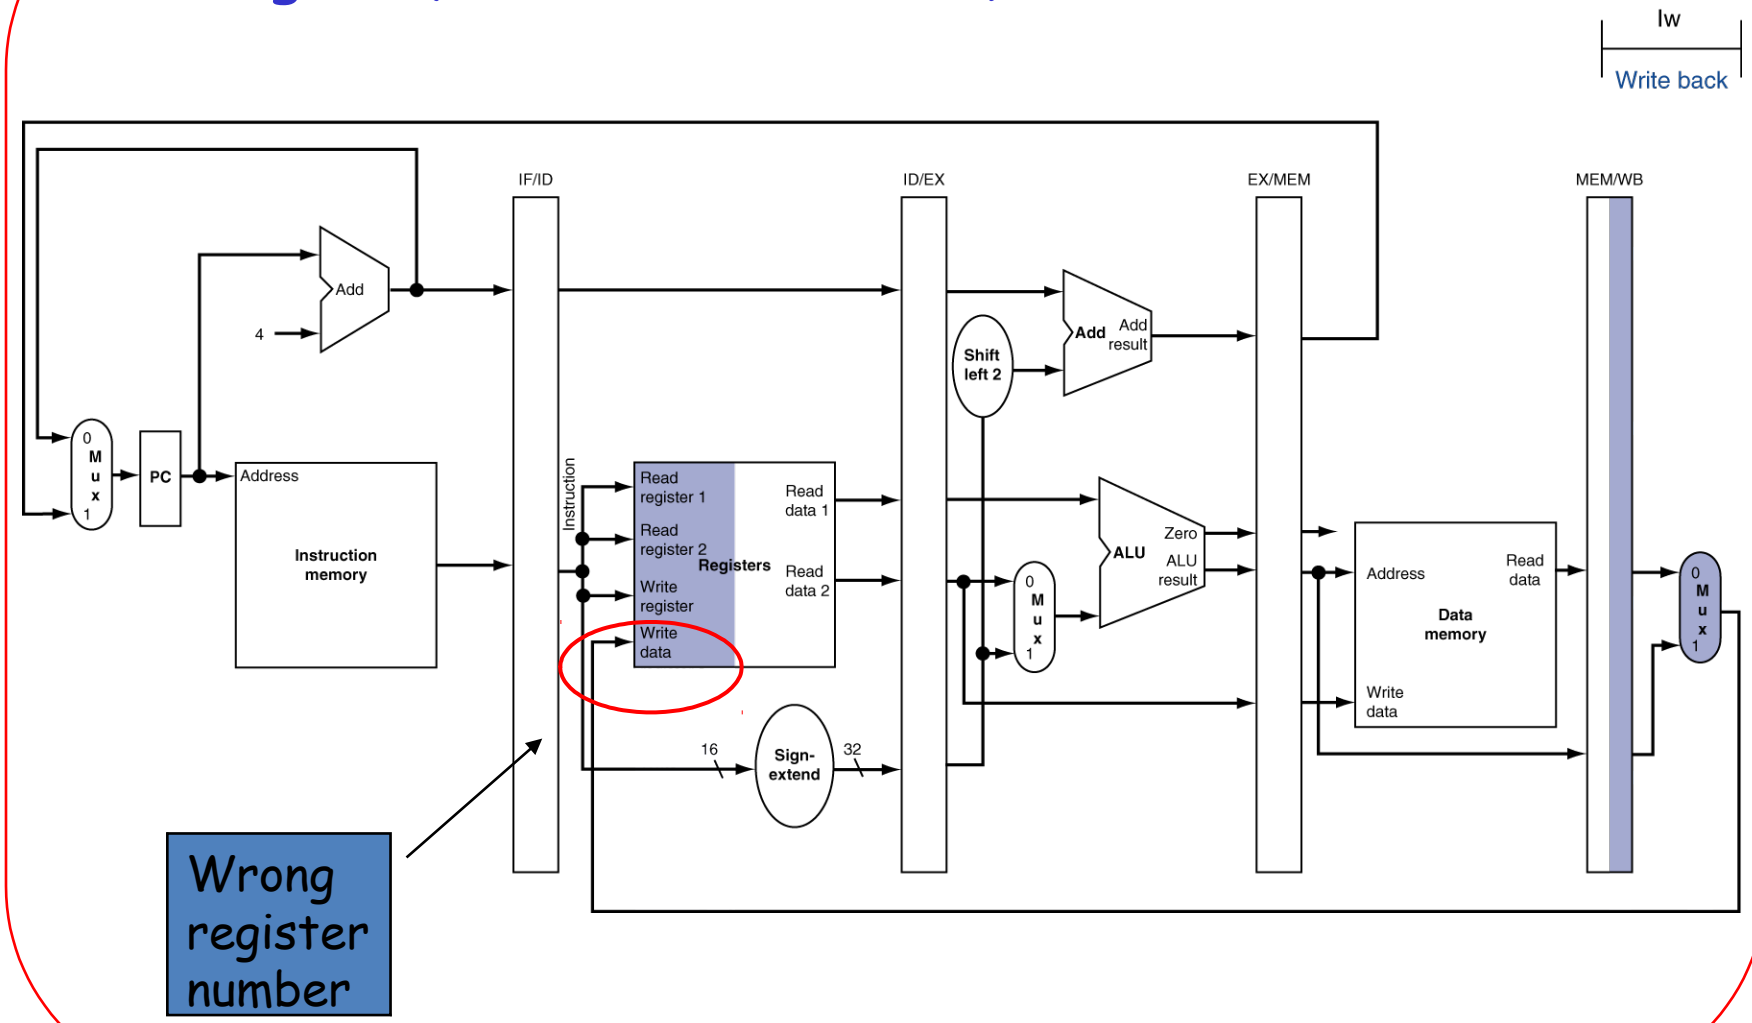

The write register has been changed

- As soon as we decode a new instruction, we will overwrite the contents of the Read Register 1 and 2 and the write register
- Hence, we have to pass along the register address to write through pipeline registers

Corrected Datapath for Load

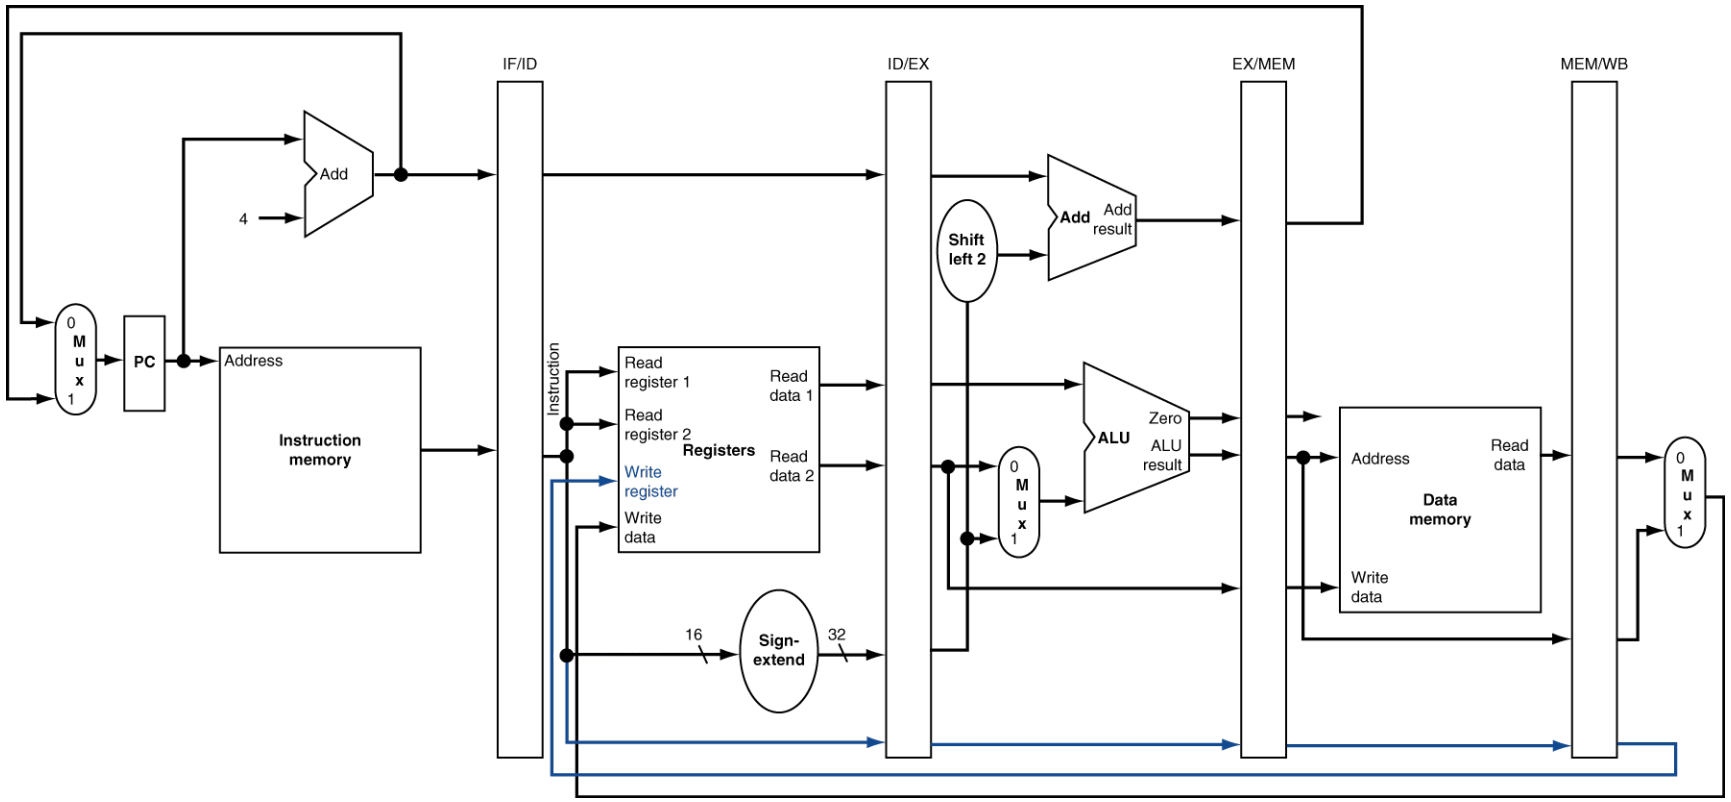

Corrected Datapath for Load

Control for the pipelined datapath

- Mostly same as for single-cycle datapath
  - Control signals are needed at corresponding stage of execution
  - No new control signals required
- But control for stage x depends on instruction at stage x
  - Control for x is decided in ID
  - Control signal is passed to pipeline registers with the instruction

Passing control down the pipeline

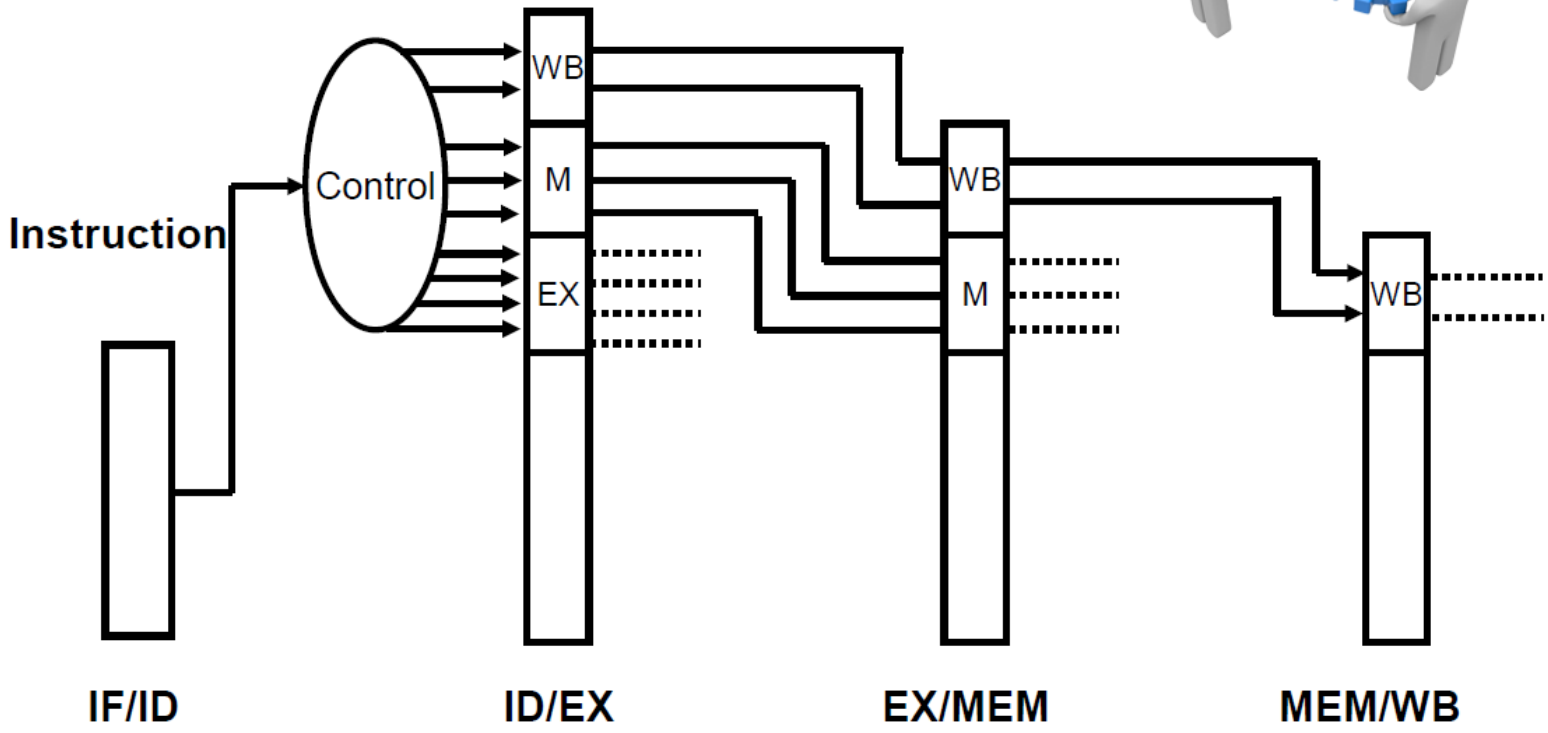

Passing control down the pipeline (full version)

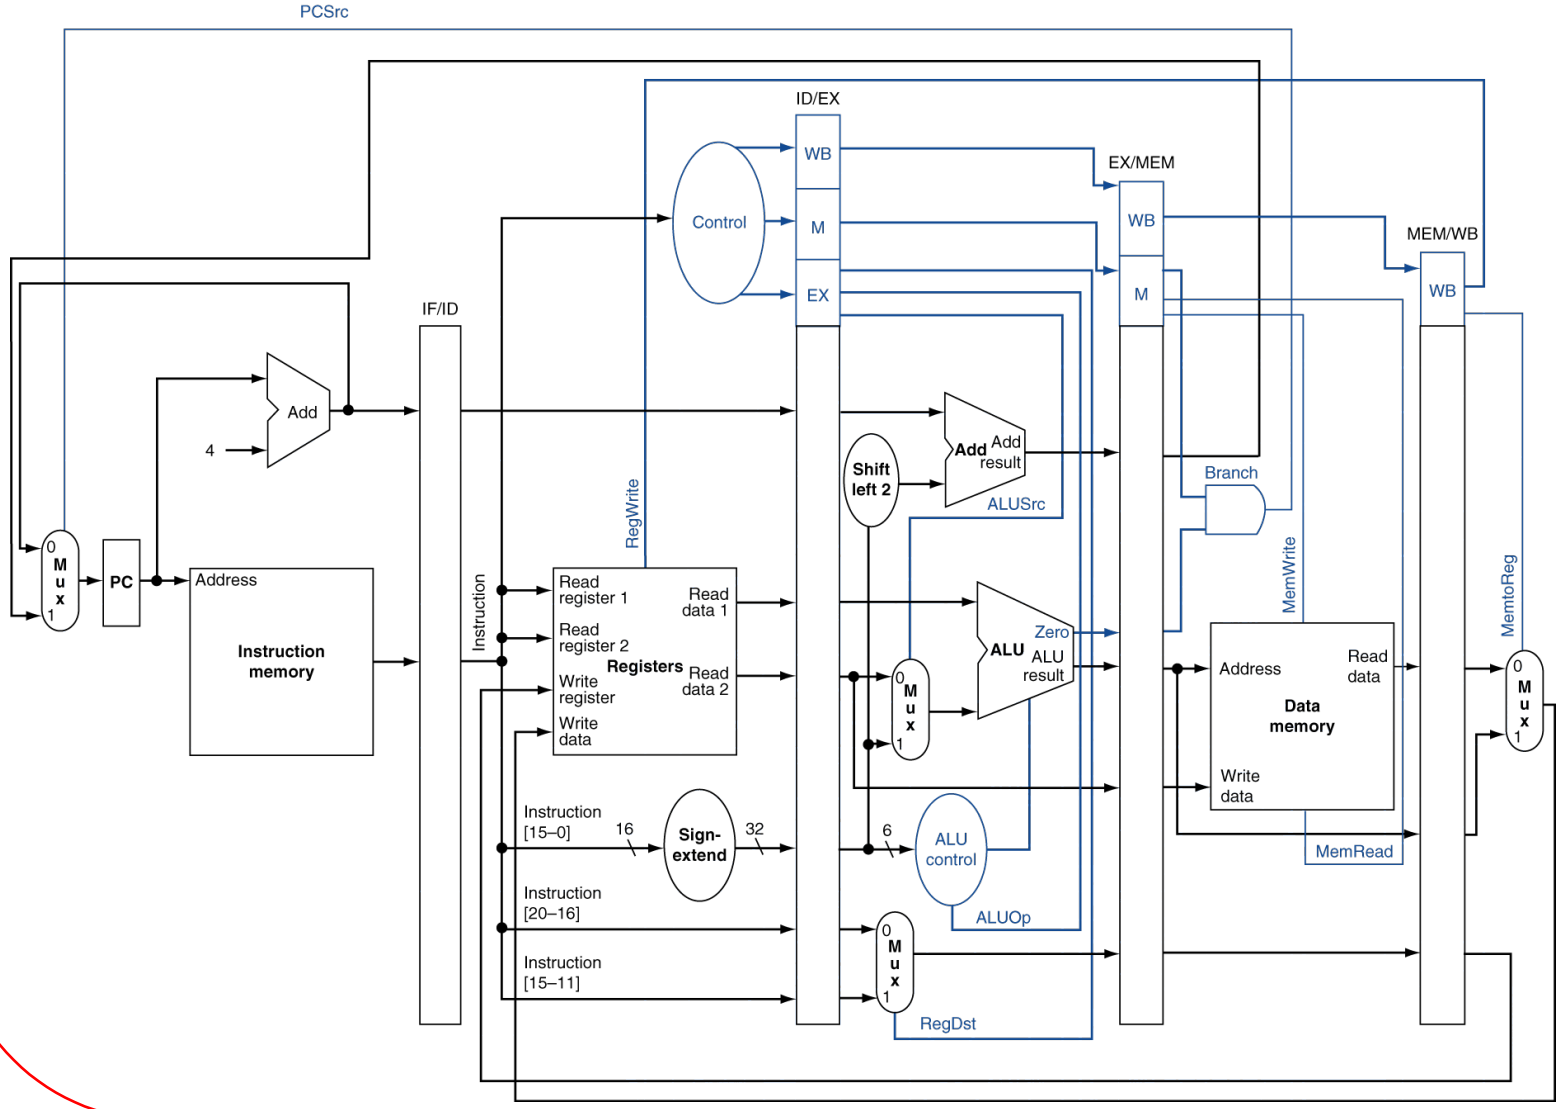

Summary

- Pipelined processors are generally faster than single-cycled processors
  - Provided the pipeline is kept full
  - The latency of an individual instruction may be greater
- Pipeline registers
  - isolate stages
  - hold information produced in previous cycle
  - get written for every cycle
  - and control hand-over of partially completed instructions
- Pipeline operation
- Control signals (values)
  - are determined in ID
  - pass through the pipeline registers as the instruction moves down the pipeline
  - are grouped and used by pipeline stage